

Kubernetes

An Introduction



Kubernetes

Pre-Requisites



Docker

YAML

Virtual Box to Setting up Lab environment,AWS
Linux



\$ whoami - Aneesh



Aneesh Ansari

er.aneeshansari@gmail.com

Sr. Devops Eng. @ MNC

Part time Freelancer

Trainer in following technology

Linux, Docker, Kubernetes, DevOps

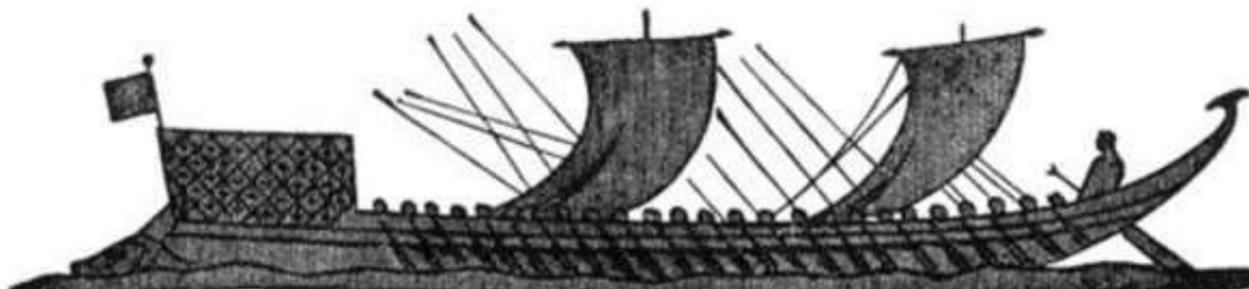


Overview

What Does “Kubernetes” Mean?



Greek for “pilot” or
“Helmsman of a ship”



[Image Source](#)



What is Kubernetes?



- Kubernetes is an open-source Container Management tool which automates container deployment, container (de)scaling & container load balancing.
- Written on Golang, it has a huge community because it was first developed by Google & later donated to **CNCF**
- Can group ‘n’ no of containers into one logical unit for managing & deploying them
- Built from the lessons learned in the experiences of developing and running Google’s Borg and Omega.

Decouples Infrastructure and Scaling



- **All services** within Kubernetes are natively Load Balanced.
- Can scale up and down dynamically.
- Used both to enable self-healing and seamless upgrading or rollback of applications.

Self Healing



Kubernetes will **ALWAYS** try and steer the cluster to its desired state.

- **Me:** “I want 3 healthy instances of redis to always be running.”
- **Kubernetes:** “Okay, I’ll ensure there are always 3 instances up and running.”
- **Kubernetes:** “Oh look, one has died. I’m going to attempt to spin up a new one.”

Most Importantly...



Use the **SAME API**
across bare metal and
EVERY cloud provider!!!



Who “Manages” Kubernetes?



The CNCF is a child entity of the Linux Foundation and operates as a vendor neutral governance group.

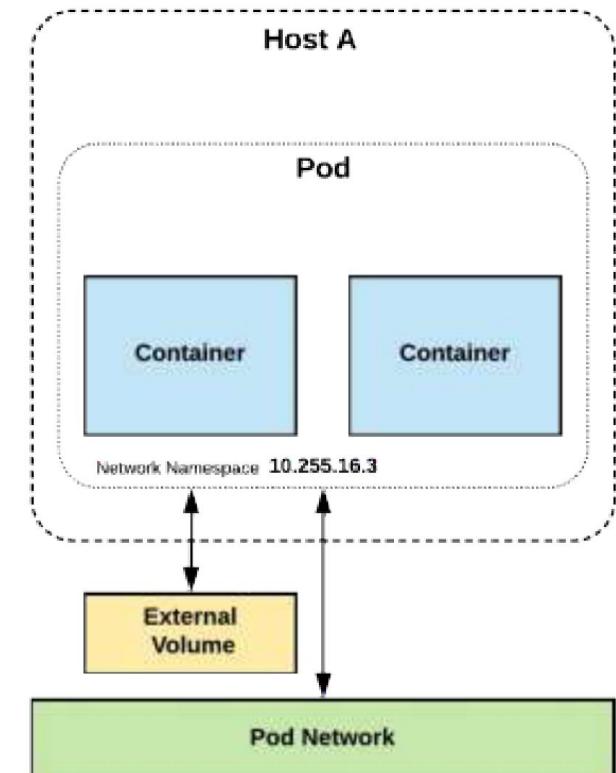


A Couple
Key Concepts...

Pods



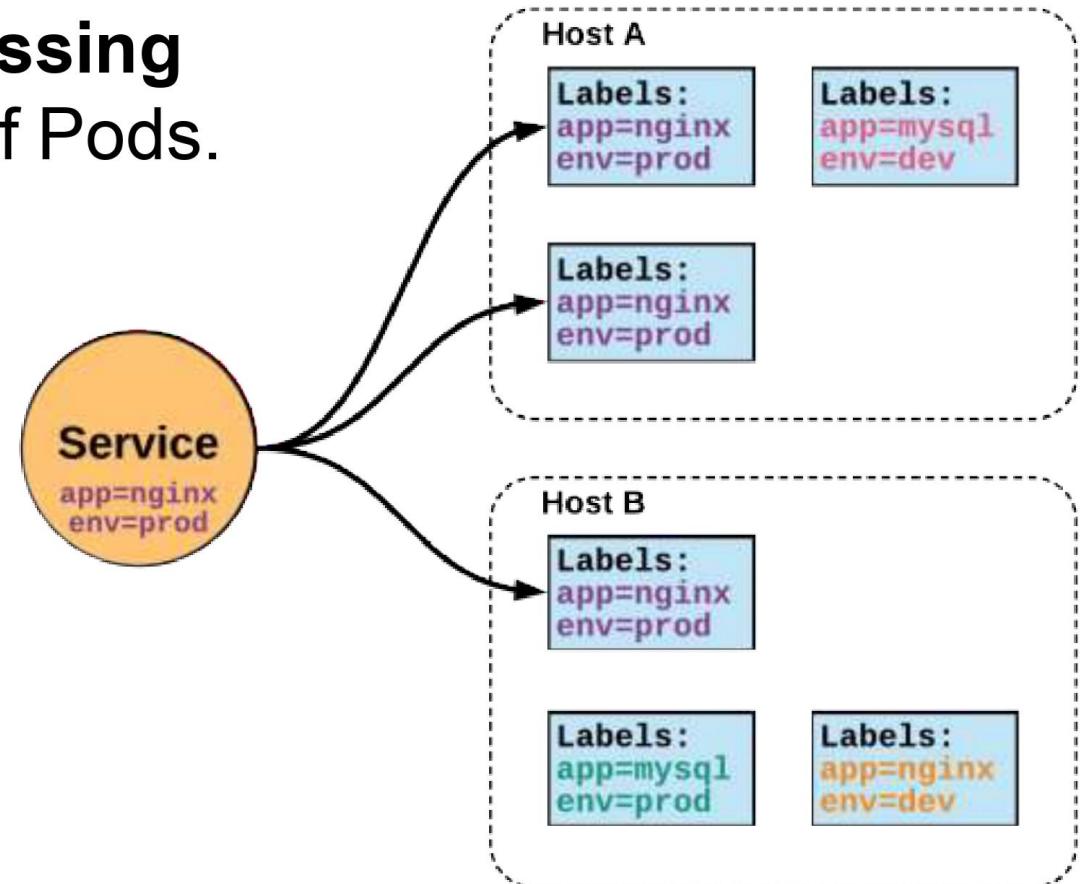
- **Atomic unit** or smallest “*unit of work*” of Kubernetes.
- Pods are **one or MORE containers** that share volumes, a network namespace, and are a part of a **single context**.



Services

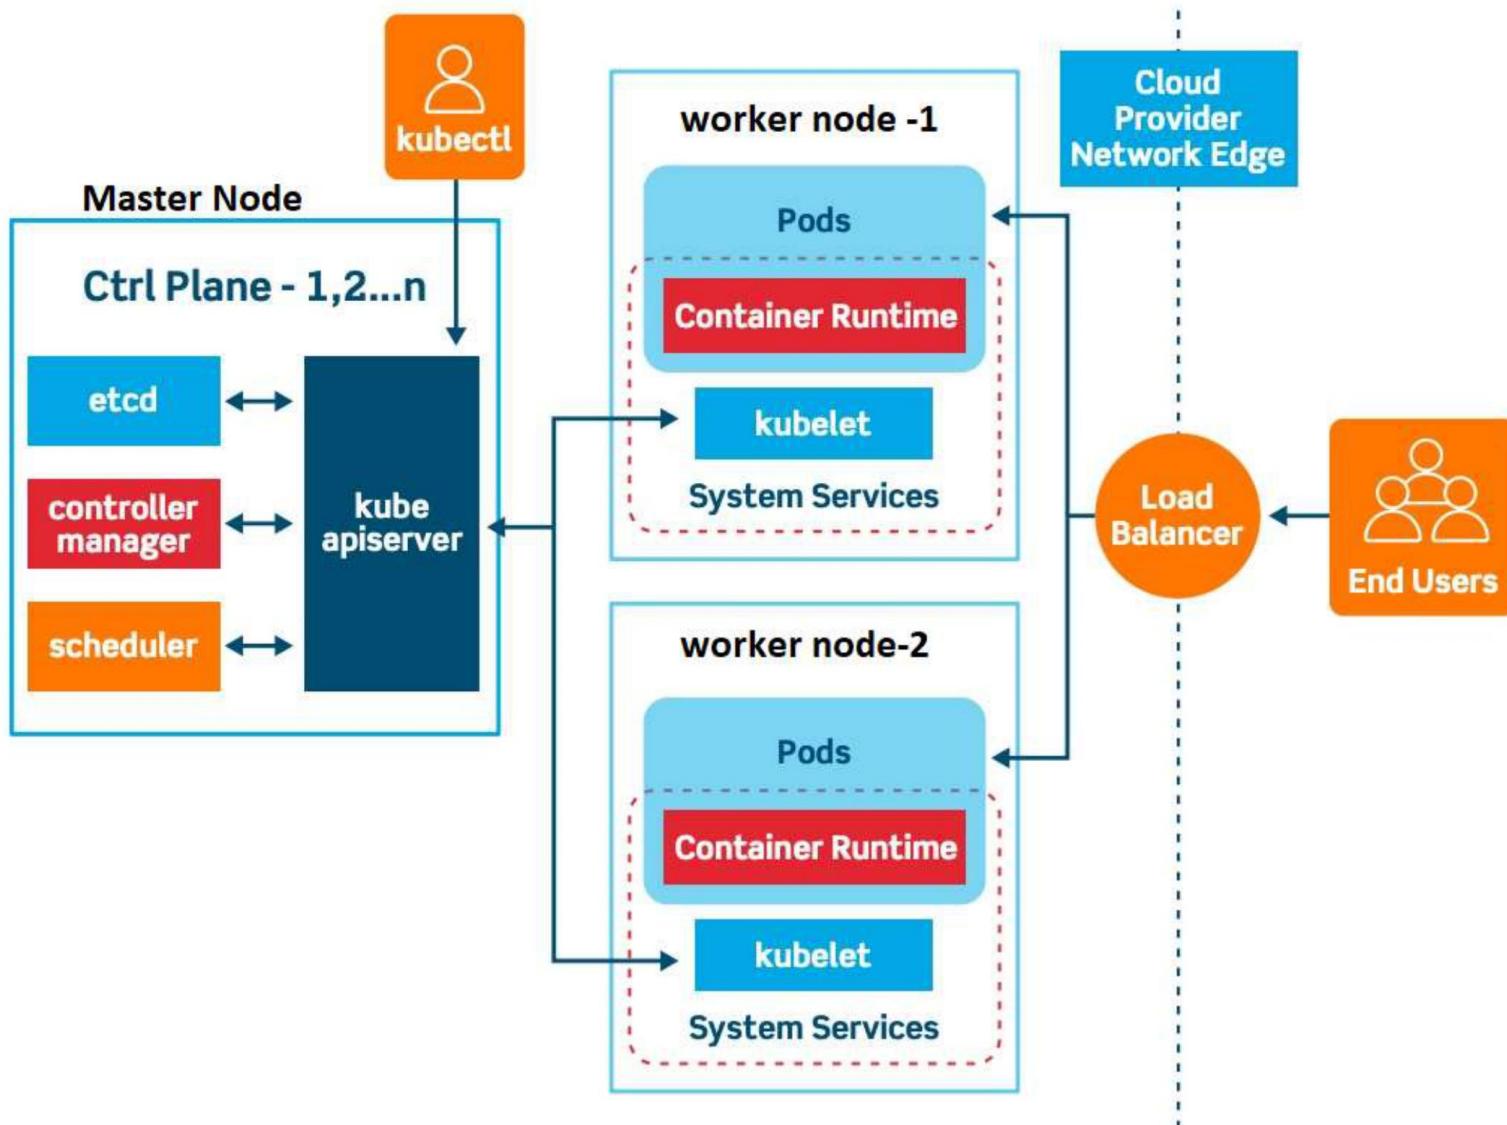


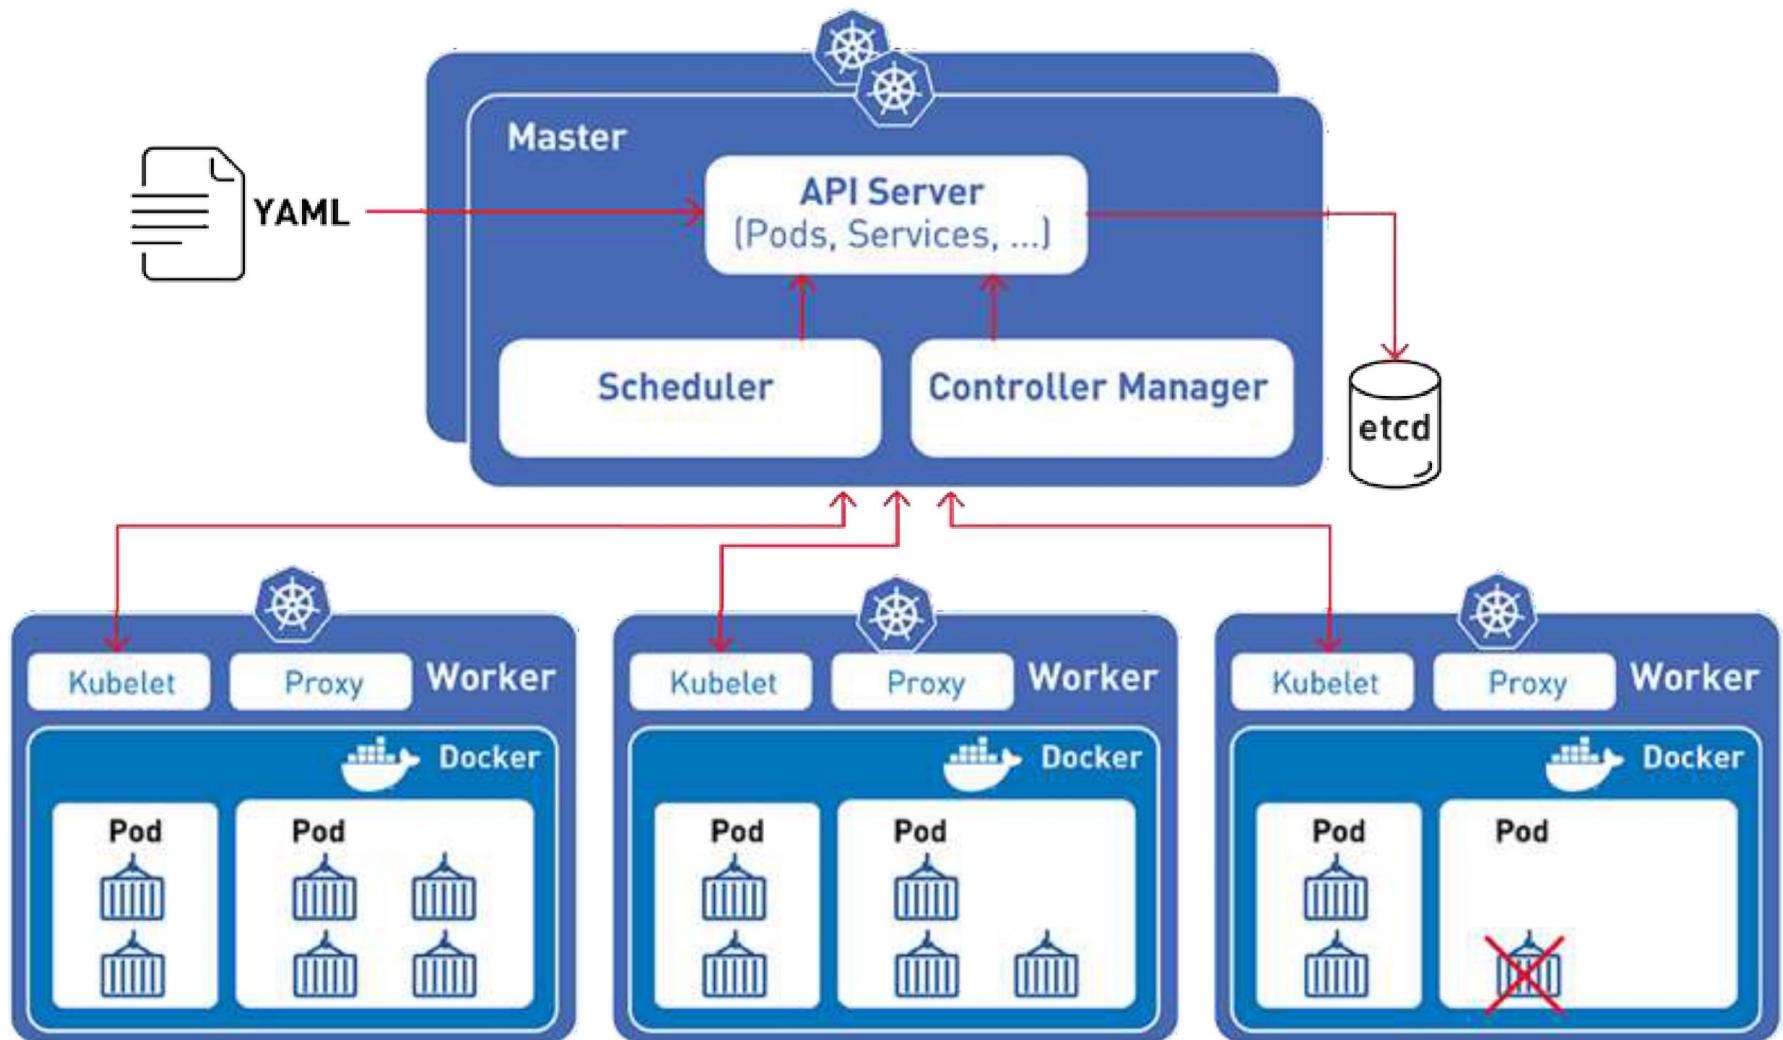
- **Unified method of accessing** the exposed workloads of Pods.
- **Durable resource**
 - static cluster IP
 - static namespaced DNS name



Architecture Overview

2379





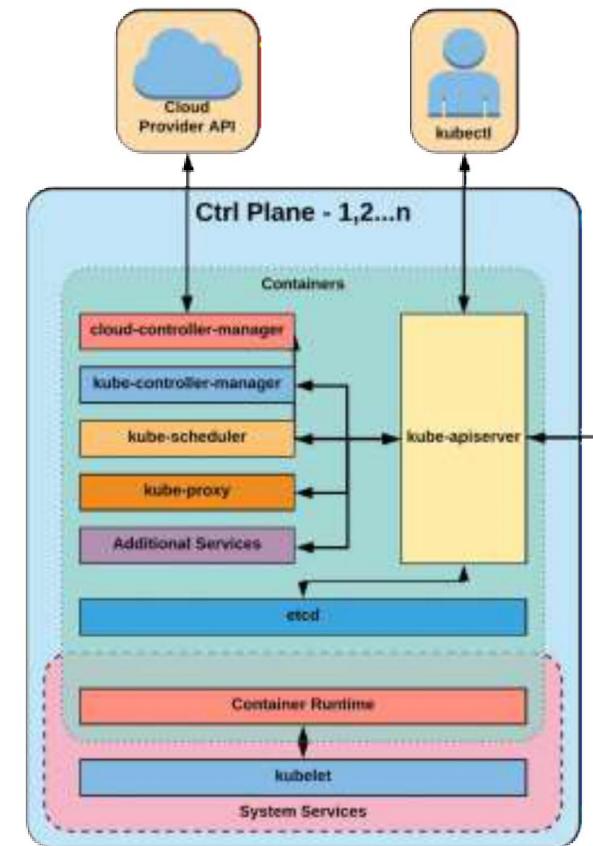
Control Plane Components

Architecture Overview

Control Plane Components



- kube-apiserver
- etcd
- kube-controller-manager
- kube-scheduler

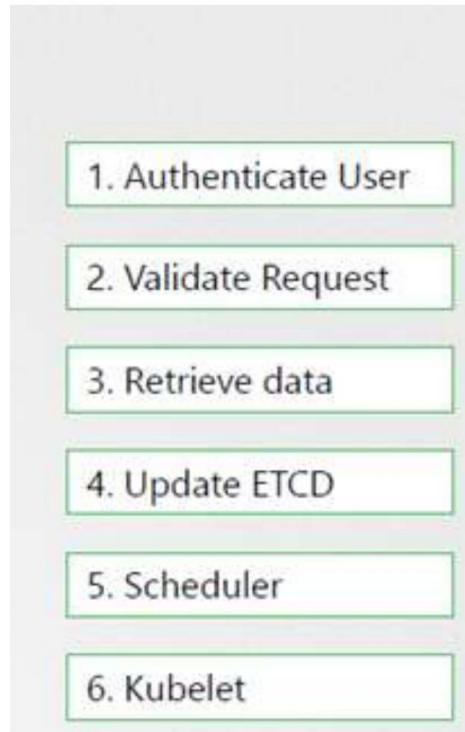


kube-apiserver



- Provides a forward facing REST interface into the kubernetes control plane and datastore.
- All clients and other applications interact with kubernetes **strictly** through the API Server.
- Acts as the gatekeeper to the cluster by handling authentication and authorization, request validation, mutation, and admission control in addition to being the front-end to the backing datastore.

kube-apiserver



etcd



- etcd acts as the cluster datastore.
- Purpose in relation to Kubernetes is to provide a strong, consistent and highly available key-value store for persisting cluster state.
- Stores objects and config information.
- If you setup cluster via kubeadm , than kubeadm will set a etcd in pod in kube-system namespace.

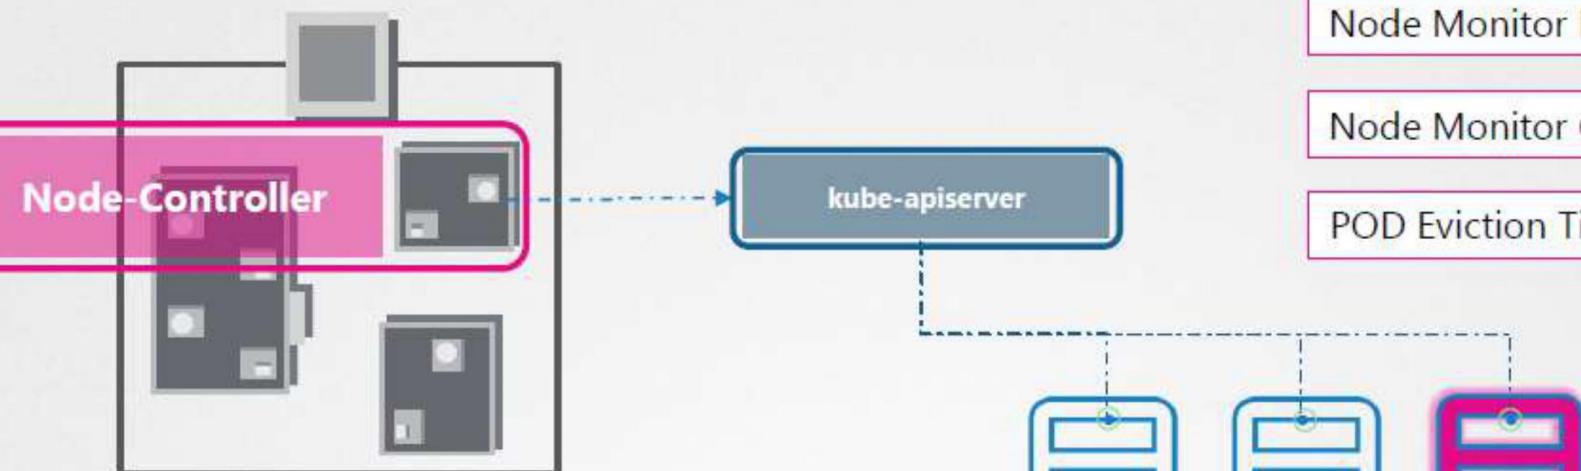


kube-controller-manager



- Serves as the primary daemon that manages all core component control loops.
- Monitors the cluster state via the apiserver and **steers the cluster towards the desired state**.
- Node and Replication Controller

Node Controller



Watch Status

Remediate Situation

Node Monitor Period = 5s

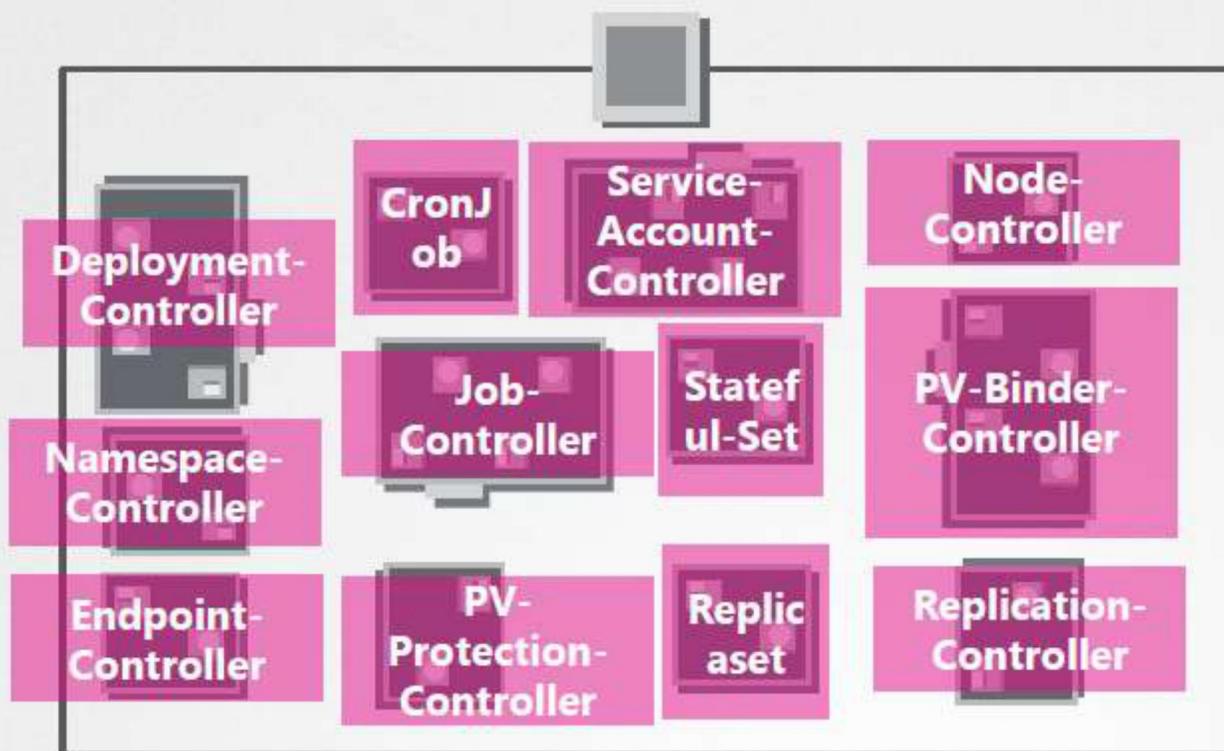
Node Monitor Grace Period = 40s

POD Eviction Timeout = 5m

```
kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
worker-1	Ready	<none>	8d	v1.13.0
worker-2	NotReady	<none>	8d	v1.13.0

IController



kube-scheduler



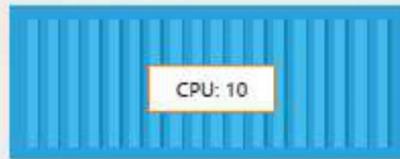
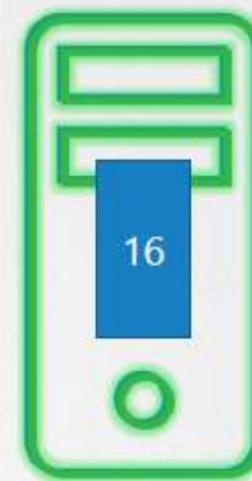
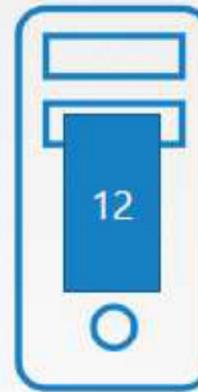
- Verbose policy-rich engine that evaluates workload requirements and attempts to place it on a matching resource.
- Scheduler only decides where pod will be deployed, than kubelet creates pod
- Workload Requirements can include: general hardware requirements, affinity/anti-affinity, labels, and other various custom resource requirements.

Kube-Scheduler



1. Filter Nodes

2. Rank Nodes



View kube-scheduler options -kubeadm

```
#cat /etc/kubernetes/manifests/kube-scheduler.yaml  
#ps-aux | grep kube-scheduler
```



```
el describe svc hello-world  
lo-world  
ault
```

KUBE-PROXY MODES

```
=hello-world  
ne>IPTABLES (default) and IP Virtual Server (IPVS).
```

```
=hello-world
```

```
ePort IPVS is an advanced configuration used in a  
ne> Cluster with thousands of services.
```

```
106.219.51
```

```
106.219.51
```

```
set> 8080/TCP
```

```
TCP
```

IPVS offers high network performance

kube-dns/proxy



Here's how Kubernetes services work! A service is a collection of pods, which each have their own IP address (like 10.1.0.3, 10.2.3.5, 10.3.5.6)

1. Every Kubernetes service gets an IP address (like 10.23.1.2)
2. `kube-dns` resolves Kubernetes service DNS names to IP addresses (so `my-svc.my-namespace.svc.cluster.local` might map to 10.23.1.2)
3. `kube-proxy` sets up iptables rules in order to do random load balancing between them. Kube-proxy also has a userspace round-robin load balancer but my impression is that they don't recommend using it.

So when you make a request to `my-svc.my-namespace.svc.cluster.local`, it resolves to 10.23.1.2, and then iptables rules on your local host (generated by kube-proxy) redirect it to one of 10.1.0.3 or 10.2.3.5 or 10.3.5.6 at random.